

# COMP2610 – Information Theory

## Lecture 11: Entropy and Coding

Robert C. Williamson

Research School of Computer Science



Australian  
National  
University

27 August 2018

# Brief Recap of Course (Last 6 Weeks)

- How can we quantify information?
  - ▶ Basic Definitions and Key Concepts
  - ▶ Probability, Entropy & Information
- How can we make good guesses?
  - ▶ Probabilistic Inference
  - ▶ Bayes Theorem
- How much redundancy can we safely remove?
  - ▶ Compression
  - ▶ Source Coding Theorem, Kraft Inequality
  - ▶ Block, Huffman, and Lempel-Ziv Coding
- How much noise can we correct and how?
  - ▶ Noisy-Channel Coding
  - ▶ Repetition Codes, Hamming Codes
- What is randomness?
  - ▶ Kolmogorov Complexity
  - ▶ Algorithmic Information Theory

# Brief Overview of Course (Next 6 Weeks)

- How can we quantify information?
  - ▶ Basic Definitions and Key Concepts
  - ▶ Probability, Entropy & Information
- How can we make good guesses?
  - ▶ Probabilistic Inference
  - ▶ Bayes Theorem
- How much redundancy can we safely remove?
  - ▶ Compression
  - ▶ Source Coding Theorem, Kraft Inequality
  - ▶ Block, Huffman, and Lempel-Ziv Coding
- How much noise can we correct and how?
  - ▶ Noisy-Channel Coding
  - ▶ Repetition Codes, Hamming Codes
- What is randomness?
  - ▶ Kolmogorov Complexity
  - ▶ Algorithmic Information Theory

# Brief Overview of Course (Next 6 Weeks)

- How can we quantify information?
  - ▶ Basic Definitions and Key Concepts
  - ▶ Probability, Entropy & Information
- How can we make good guesses?
  - ▶ Probabilistic Inference
  - ▶ Bayes Theorem
- How much redundancy can we safely remove?
  - ▶ Compression
  - ▶ Source Coding Theorem, Kraft Inequality
  - ▶ Block, Huffman, and Lempel-Ziv Coding
- How much noise can we correct and how?
  - ▶ Noisy-Channel Coding
  - ▶ Repetition Codes, Hamming Codes
- What is randomness?
  - ▶ Kolmogorov Complexity
  - ▶ Algorithmic Information Theory

## This time

Basic goal of compression

Key concepts: codes and their types, raw bit content, essential bit content

Informal statement of source coding theorem

## 1 Introduction

- Overview
- What is Compression?
- A Communication Game
- What's the best we can do?

## 2 Formalising Coding

- Entropy and Information: A Quick Review
- Defining Codes

## 3 Formalising Compression

- Reliability vs. Size
- Key Result: The Source Coding Theorem


# What is Compression?

Cn y rd ths mssg wtht ny vwls?

# What is Compression?

Cn y rd ths mssg wtht ny vwls?

It is not too difficult to read as there is **redundancy** in English text.  
(Estimates of 1-1.5 bits per character, compared to  $\log_2 26 \approx 4.7$ )



(a)  $P(y|x)$

- If you see a “q”, it is very likely to be followed with a “u”
- The letter “e” is much more common than “j”
- Compression exploits differences in relative probability of symbols or blocks of symbols



# Compression in a Nutshell


## Compression

Data compression is the process of replacing a message with a **smaller message** which can be **reliably converted** back to the original.

# A General Communication Game

Imagine the following game between Sender & Receiver:

- Sender & Receiver agree on **code** for each outcome ahead of time (e.g., 0 for *Heads*; 1 for *Tails*)
- Sender observes outcomes then codes and sends message
- Receiver decodes message and recovers outcome sequence



**Goal:** Want small messages **on average** when outcomes are from a **fixed, known, but uncertain** source (e.g., coin flips with known bias)

## Sneak peek: source coding theorem

Consider a coin with  $P(\text{Heads}) = 0.9$ . If we want perfect transmission:

- Coding single outcomes requires 1 bit/outcome
- Coding 10 outcomes at a time needs 10 bits, or 1 bit/outcome

Not very interesting!

## Sneak peek: source coding theorem

Consider a coin with  $P(\text{Heads}) = 0.9$ . If we want perfect transmission:

- Coding single outcomes requires 1 bit/outcome
- Coding 10 outcomes at a time needs 10 bits, or 1 bit/outcome

Not very interesting!

Things get interesting if we:

- accept errors in transmission
- allow variable length messages

# Sneak peek: source coding theorem

Consider a coin with  $P(\text{Heads}) = 0.9$ . If we want perfect transmission:

- Coding single outcomes requires 1 bit/outcome
- Coding 10 outcomes at a time needs 10 bits, or 1 bit/outcome

Not very interesting!

Things get interesting if we:

- **accept errors in transmission** (this week)
- allow variable length messages (next week)

## Sneak peek: source coding theorem

If we are happy to fail on up to 2% of the sequences we can ignore any sequence of 10 outcomes with more than 3 tails

*Why?* The number of tails follows a Binomial(10, 0.1) distribution

## Sneak peek: source coding theorem

If we are happy to fail on up to 2% of the sequences we can ignore any sequence of 10 outcomes with more than **3 tails**

*Why?* The number of tails follows a Binomial(10, 0.1) distribution

There are only  $176 < 2^8$  sequences with 3 or fewer tails

So, we can just code those, and **ignore** the rest!

- Coding 10 outcomes with 2% failure doable with 8 bits, or **0.8 bits/outcome**
- *Smallest bits/outcome needed for 10,000 outcome sequences?*

# Generalisation: Source Coding Theorem

What happens when we generalise to arbitrary error probability, and sequence size?



# Generalisation: Source Coding Theorem

What happens when we generalise to arbitrary error probability, and sequence size?

## Source Coding Theorem (Informal Statement)

**If:** you want to uniformly code large sequences of outcomes with any degree of reliability from a random source

**Then:** the average number of bits per outcome you will **need** is roughly equal to the entropy of that source.

# Generalisation: Source Coding Theorem

What happens when we generalise to arbitrary error probability, and sequence size?

## Source Coding Theorem (Informal Statement)

**If:** you want to **uniformly code large sequences** of outcomes with any **degree of reliability** from a random source

**Then:** the **average number of bits per outcome** you will **need** is **roughly equal** to the entropy of that source.

**To define:** “Uniformly code”, “large sequences”, “degree of reliability”, “average number of bits per outcome”, “roughly equal”

## 1 Introduction

- Overview
- What is Compression?
- A Communication Game
- What's the best we can do?

## 2 Formalising Coding

- Entropy and Information: A Quick Review
- Defining Codes

## 3 Formalising Compression

- Reliability vs. Size
- Key Result: The Source Coding Theorem

# Entropy and Information: Recap

## Ensemble

An **ensemble**  $X$  is a triple  $(x, \mathcal{A}_X, \mathcal{P}_X)$ ;  $x$  is a **random variable** taking **values** in  $\mathcal{A}_X = \{a_1, a_2, \dots, a_I\}$  with **probabilities**  $\mathcal{P}_X = \{p_1, p_2, \dots, p_I\}$ .

# Entropy and Information: Recap

## Ensemble

An **ensemble**  $X$  is a triple  $(x, \mathcal{A}_X, \mathcal{P}_X)$ ;  $x$  is a **random variable** taking **values** in  $\mathcal{A}_X = \{a_1, a_2, \dots, a_I\}$  with **probabilities**  $\mathcal{P}_X = \{p_1, p_2, \dots, p_I\}$ .

## Information

The **information** in the observation that  $x = a_i$  (in the ensemble  $X$ ) is

$$h(a_i) = \log_2 \frac{1}{p_i} = -\log_2 p_i$$

# Entropy and Information: Recap

## Ensemble

An **ensemble**  $X$  is a triple  $(x, \mathcal{A}_X, \mathcal{P}_X)$ ;  $x$  is a **random variable** taking **values** in  $\mathcal{A}_X = \{a_1, a_2, \dots, a_I\}$  with **probabilities**  $\mathcal{P}_X = \{p_1, p_2, \dots, p_I\}$ .

## Information

The **information** in the observation that  $x = a_i$  (in the ensemble  $X$ ) is

$$h(a_i) = \log_2 \frac{1}{p_i} = -\log_2 p_i$$

## Entropy

The **entropy** of an ensemble  $X$  is the average information

$$H(X) = \mathbb{E}[h(x)] = \sum_i p_i h(a_i) = \sum_i p_i \log_2 \frac{1}{p_i}$$

# What is a Code?

A source code is a process for assigning **names** to outcomes. The names are typically expressed by **strings of binary symbols**.

We will denote the set of all finite binary strings by

$$\{0, 1\}^+ \stackrel{\text{def}}{=} \{0, 1, 00, 01, 10, \dots\}$$

## Source Code

Given an ensemble  $X$ , the function  $c : \mathcal{A}_X \rightarrow \{0, 1\}^+$  is a **source code** for  $X$ . The number of symbols in  $c(x)$  is the **length**  $l(x)$  of the codeword for  $x$ . The **extension** of  $c$  is defined by  $c(x_1 \dots x_n) = c(x_1) \dots c(x_n)$

# What is a Code?

A source code is a process for assigning **names** to outcomes. The names are typically expressed by **strings of binary symbols**.

We will denote the set of all finite binary strings by

$$\{0, 1\}^+ \stackrel{\text{def}}{=} \{0, 1, 00, 01, 10, \dots\}$$

## Source Code

Given an ensemble  $X$ , the function  $c : \mathcal{A}_X \rightarrow \{0, 1\}^+$  is a **source code** for  $X$ . The number of symbols in  $c(x)$  is the **length**  $l(x)$  of the codeword for  $x$ . The **extension** of  $c$  is defined by  $c(x_1 \dots x_n) = c(x_1) \dots c(x_n)$

### Example:

- The **code**  $c$  names outcomes from  $\mathcal{A}_X = \{\mathbf{r}, \mathbf{g}, \mathbf{b}\}$  by  $c(\mathbf{r}) = 00$ ,  $c(\mathbf{g}) = 10$ ,  $c(\mathbf{b}) = 11$
- The **length** of the codeword for each outcome is 2.
- The **extension** of  $c$  gives  $c(\mathbf{rgrb}) = 00100011$



# Types of Codes

Let  $X$  be an ensemble and  $c : \mathcal{A}_X \rightarrow \{0, 1\}^+$  a code for  $X$ . We say  $c$  is a:

- **Uniform Code** if  $l(x)$  is the same for all  $x \in \mathcal{A}_X$
- **Variable-Length Code** otherwise

# Types of Codes

Let  $X$  be an ensemble and  $c : \mathcal{A}_X \rightarrow \{0, 1\}^+$  a code for  $X$ . We say  $c$  is a:

- **Uniform Code** if  $l(x)$  is the same for all  $x \in \mathcal{A}_X$
- **Variable-Length Code** otherwise

Another important criteria for codes is whether the original symbol  $x$  can be unambiguously determined given  $c(x)$ . We say  $c$  is a:

- **Lossless Code** if for all  $x_1, x_2 \in \mathcal{A}_X$  we have  $x_1 \neq x_2$  implies  $c(x_1) \neq c(x_2)$
- **Lossy Code** otherwise

# Types of Codes

## Examples

**Examples:** Let  $\mathcal{A}_X = \{a, b, c, d\}$

- ➊  $c(a) = 00, c(b) = 01, c(c) = 10, c(d) = 11$  is uniform lossless
- ➋  $c(a) = 0, c(b) = 10, c(c) = 110, c(d) = 111$  is variable-length lossless
- ➌  $c(a) = 0, c(b) = 0, c(c) = 110, c(d) = 111$  is variable-length lossy
- ➍  $c(a) = 00, c(b) = 00, c(c) = 10, c(d) = 11$  is uniform lossy
- ➎  $c(a) = -, c(b) = -, c(c) = 10, c(d) = 11$  is uniform lossy

## A Note on Lossy Codes & Missing Codewords

When talking about a uniform lossy code  $c$  for  $\mathcal{A}_X = \{a, b, c\}$  we write

$$c(a) = 0 \quad c(b) = 1 \quad c(c) = -$$

where the symbol  $-$  means “no codeword”. This is shorthand for “the receiver will decode this codeword incorrectly”

For the purposes of these lectures, this is equivalent to the code

$$c(a) = 0 \quad c(b) = 1 \quad c(c) = 1$$

and the sender and receiver agreeing that the codeword 1 should always be decoded as  $b$

Assigning the outcome  $a_i$  the missing codeword “ $-$ ” just means “it is not possible to send  $a_i$  reliably”

## 1 Introduction

- Overview
- What is Compression?
- A Communication Game
- What's the best we can do?

## 2 Formalising Coding

- Entropy and Information: A Quick Review
- Defining Codes

## 3 Formalising Compression

- Reliability vs. Size
- Key Result: The Source Coding Theorem

# Lossless Coding

Example: Colours



Three colour ensemble with  $\mathcal{A}_X = \{\text{r}, \text{g}, \text{b}\}$   
with **r** twice as likely as **b** or **g**

- $p_{\text{r}} = 0.5$  and  $p_{\text{g}} = p_{\text{b}} = 0.25$ .

Suppose we use the following **uniform lossless** code

$$c(\text{r}) = 00; c(\text{g}) = 10; \text{ and } c(\text{b}) = 11$$

For example  $c(\text{rrgbrbr}) = 00001011001100$  will have 14 bits.

On average, we will use  $I(\text{r})p_{\text{r}} + I(\text{g})p_{\text{g}} + I(\text{b})p_{\text{b}} = 2$  bits per outcome

- $2N$  bits to code a sequence of  $N$  outcomes

# Raw Bit Content

Uniform coding gives a crude measure of information: the **number of bits required to assign equal length codes to each symbol**

## Raw Bit Content

If  $X$  is an ensemble with outcome set  $\mathcal{A}_X$  then its **raw bit content** is

$$H_0(X) = \log_2 |\mathcal{A}_X|.$$

$x$	$c(x)$
a	000
b	001
c	010
d	011
e	100
f	101
g	110
h	111

### Example:

This is a uniform encoding of outcomes in

$\mathcal{A}_X = \{a, b, c, d, e, f, g, h\}$ :

- Each outcome is encoded using  $H_0(X) = 3$  bits
- The **probabilities** of the outcomes **are ignored**
- Same as assuming a **uniform distribution**

For the purposes of compression, the exact codes don't matter – just the number of bits used.

# Lossy Coding

Example: Colours



Three colour ensemble with  $\mathcal{A}_X = \{\mathbf{r}, \mathbf{g}, \mathbf{b}\}$

- $p_{\mathbf{r}} = 0.5$  and  $p_{\mathbf{g}} = p_{\mathbf{b}} = 0.25$ .

Using **uniform lossy** code:

- $c(\mathbf{r}) = 0$ ;  $c(\mathbf{g}) = -$ ; and  $c(\mathbf{b}) = 1$

**Examples:**

$c(\mathbf{rrrrrrrr}) = 0000000$ ;  $c(\mathbf{rrbbrbr}) = 0011010$ ;  $c(\mathbf{rrgbrbr}) = -$



# Lossy Coding

Example: Colours



Three colour ensemble with  $\mathcal{A}_X = \{\mathbf{r}, \mathbf{g}, \mathbf{b}\}$

- $p_{\mathbf{r}} = 0.5$  and  $p_{\mathbf{g}} = p_{\mathbf{b}} = 0.25$ .

Using **uniform lossy** code:

- $c(\mathbf{r}) = 0$ ;  $c(\mathbf{g}) = -$ ; and  $c(\mathbf{b}) = 1$

**Examples:**

$c(\mathbf{rrrrrrrr}) = 0000000$ ;  $c(\mathbf{rrbbrbr}) = 0011010$ ;  $c(\mathbf{rrgbrbr}) = -$

What is **probability we can reliably code** a sequence of  $N$  outcomes?

Given we can code a sequence of length  $N$ , **how many bits are expected?**

# Lossy Coding

Example: Colours

What is **probability we can reliably code** a sequence of  $N$  outcomes?

$$P(x_1 \dots x_N \text{ has no } g) = P(x_1 \neq g) \dots P(x_N \neq g) = (1 - p_g)^N$$

# Lossy Coding

Example: Colours

What is **probability we can reliably code** a sequence of  $N$  outcomes?

$$P(x_1 \dots x_N \text{ has no } g) = P(x_1 \neq g) \dots P(x_N \neq g) = (1 - p_g)^N$$

Given we can code a sequence of length  $N$ , **how many bits are expected?**

$$\begin{aligned} \mathbb{E}[I(X_1) + \dots + I(X_N) | X_1 \neq g, \dots, X_N \neq g] &= \sum_{n=1}^N \mathbb{E}[I(X_n) | X_n \neq g] \\ &= N(I(r)p_r + I(b)p_b) / (1 - p_g) = N \end{aligned}$$

since  $I(p_r) = I(p_b) = 1$  and  $p_r + p_b = 1 - p_g$ .

- c.f.  $2N$  bits with lossless code

# Lossy Coding

Example: Colours

What is **probability we can reliably code** a sequence of  $N$  outcomes?

$$P(x_1 \dots x_N \text{ has no } g) = P(x_1 \neq g) \dots P(x_N \neq g) = (1 - p_g)^N$$

Given we can code a sequence of length  $N$ , **how many bits are expected?**

$$\begin{aligned} \mathbb{E}[I(X_1) + \dots + I(X_N) | X_1 \neq g, \dots, X_N \neq g] &= \sum_{n=1}^N \mathbb{E}[I(X_n) | X_n \neq g] \\ &= N(I(r)p_r + I(b)p_b) / (1 - p_g) = N \end{aligned}$$

since  $I(p_r) = I(p_b) = 1$  and  $p_r + p_b = 1 - p_g$ .

- c.f.  $2N$  bits with lossless code

# Lossy Coding

Example: Colours

What is **probability we can reliably code** a sequence of  $N$  outcomes?

$$P(x_1 \dots x_N \text{ has no } g) = P(x_1 \neq g) \dots P(x_N \neq g) = (1 - p_g)^N$$

Given we can code a sequence of length  $N$ , **how many bits are expected?**

$$\begin{aligned} \mathbb{E}[I(X_1) + \dots + I(X_N) | X_1 \neq g, \dots, X_N \neq g] &= \sum_{n=1}^N \mathbb{E}[I(X_n) | X_n \neq g] \\ &= N(I(r)p_r + I(b)p_b) / (1 - p_g) = N = N \log_2 |\{r, b\}| \end{aligned}$$

since  $I(p_r) = I(p_b) = 1$  and  $p_r + p_b = 1 - p_g$ .

- c.f.  $2N$  bits with lossless code

# Essential Bit Content

There is an inherent trade off between the number of bits required in a uniform lossy code and the probability of being able to code an outcome

## Smallest $\delta$ -sufficient subset

Let  $X$  be an ensemble and for  $0 \leq \delta \leq 1$ , define  $S_\delta$  to be the smallest subset of  $\mathcal{A}_X$  such that

$$P(x \in S_\delta) \geq 1 - \delta$$

For small  $\delta$ , **smallest** collection of **most likely** outcomes

# Essential Bit Content

There is an inherent trade off between the number of bits required in a uniform lossy code and the probability of being able to code an outcome

## Smallest $\delta$ -sufficient subset

Let  $X$  be an ensemble and for  $0 \leq \delta \leq 1$ , define  $S_\delta$  to be the smallest subset of  $\mathcal{A}_X$  such that

$$P(x \in S_\delta) \geq 1 - \delta$$

For small  $\delta$ , **smallest** collection of **most likely** outcomes

If we uniformly code elements in  $S_\delta$ , and ignore all others:

- We can code a sequence of length  $N$  with probability  $(1 - \delta)^N$
- If we can code a sequence, its expected length is  $N \log_2 |S_\delta|$

# Essential Bit Content

## Example

Intuitively, construct  $S_\delta$  by removing elements of  $X$  in ascending order of probability, till we have reached the  $1 - \delta$  threshold

$x$	$P(x)$
a	1/4
b	1/4
c	1/4
d	3/16
e	1/64
f	1/64
g	1/64
h	1/64

- Outcomes ranked (high–low) by  $P(x = a_i)$   
removed to make set  $S_\delta$  with  $P(x \in S_\delta) \geq 1 - \delta$

$$\delta = 0 : S_\delta = \{a, b, c, d, e, f, g, h\}$$



# Essential Bit Content

## Example

Intuitively, construct  $S_\delta$  by removing elements of  $X$  in ascending order of probability, till we have reached the  $1 - \delta$  threshold

---

$x$	$P(x)$
a	1/4
b	1/4
c	1/4
d	3/16
e	1/64
f	1/64
g	1/64

---

- Outcomes ranked (high–low) by  $P(x = a_i)$   
removed to make set  $S_\delta$  with  $P(x \in S_\delta) \geq 1 - \delta$

$$\delta = 0 : S_\delta = \{a, b, c, d, e, f, g, h\}$$

$$\delta = 1/64 : S_\delta = \{a, b, c, d, e, f, g\}$$

# Essential Bit Content

## Example

Intuitively, construct  $S_\delta$  by removing elements of  $X$  in ascending order of probability, till we have reached the  $1 - \delta$  threshold

$x$	$P(x)$
a	1/4
b	1/4
c	1/4
d	3/16

- Outcomes ranked (high–low) by  $P(x = a_i)$   
removed to make set  $S_\delta$  with  $P(x \in S_\delta) \geq 1 - \delta$

$$\delta = 0 : S_\delta = \{a, b, c, d, e, f, g, h\}$$

$$\delta = 1/64 : S_\delta = \{a, b, c, d, e, f, g\}$$

$$\delta = 1/16 : S_\delta = \{a, b, c, d\}$$

# Essential Bit Content

## Example

Intuitively, construct  $S_\delta$  by removing elements of  $X$  in ascending order of probability, till we have reached the  $1 - \delta$  threshold

$\mathbf{x}$	$P(\mathbf{x})$
a	1/4

- Outcomes ranked (high–low) by  $P(x = a_i)$   
removed to make set  $S_\delta$  with  $P(x \in S_\delta) \geq 1 - \delta$

$$\delta = 0 : S_\delta = \{a, b, c, d, e, f, g, h\}$$

$$\delta = 1/64 : S_\delta = \{a, b, c, d, e, f, g\}$$

$$\delta = 1/16 : S_\delta = \{a, b, c, d\}$$

$$\delta = 3/4 : S_\delta = \{a\}$$

# Essential Bit Content

Trade off between a probability of  $\delta$  of not coding an outcome and size of uniform code is captured by the **essential bit content**

## Essential Bit Content

For an ensemble  $X$  and  $0 \leq \delta \leq 1$ , the **essential bit content** of  $X$  is

$$H_{\delta}(X) \stackrel{\text{def}}{=} \log_2 |S_{\delta}|$$

# Essential Bit Content


Trade off between a probability of  $\delta$  of not coding an outcome and size of uniform code is captured by the **essential bit content**

## Essential Bit Content

For an ensemble  $X$  and  $0 \leq \delta \leq 1$ , the **essential bit content** of  $X$  is

$$H_\delta(X) \stackrel{\text{def}}{=} \log_2 |S_\delta|$$

$\mathbf{x}$	$P(\mathbf{x})$
a	1/4
b	1/4
c	1/4
d	3/16
e	1/64
f	1/64
g	1/64
h	1/64



# The Source Coding Theorem for Uniform Codes

(Theorem 4.1 in MacKay)

Our aim next time is to understand this:

## The Source Coding Theorem for Uniform Codes

Let  $X$  be an ensemble with entropy  $H = H(X)$  bits. Given  $\epsilon > 0$  and  $0 < \delta < 1$ , there exists a positive integer  $N_0$  such that for all  $N > N_0$

$$\left| \frac{1}{N} H_\delta(X^N) - H \right| < \epsilon.$$

# The Source Coding Theorem for Uniform Codes

(Theorem 4.1 in MacKay)

Our aim next time is to understand this:

## The Source Coding Theorem for Uniform Codes

Let  $X$  be an ensemble with entropy  $H = H(X)$  bits. Given  $\epsilon > 0$  and  $0 < \delta < 1$ , there exists a positive integer  $N_0$  such that for all  $N > N_0$

$$\left| \frac{1}{N} H_{\delta}(X^N) - H \right| < \epsilon.$$

### What?

- The term  $\frac{1}{N} H_{\delta}(X^N)$  is the average number of bits required to **uniformly** code all but a proportion  $\delta$  of the symbols.
- Given a **tiny** probability of error  $\delta$ , the average bits per symbol can be made as close to  $H$  as required.
- Even if we allow a **large** probability of error we cannot compress more than  $H$  bits per symbol.

# Some Intuition for the SCT

- Don't code individual symbols in an ensemble; rather, consider sequences of length  $N$ .
- As length of sequence increases, the probability of seeing a “typical” sequence becomes much larger than “atypical” sequences.
- Thus, we can get by with essentially assigning a unique codeword to each typical sequence



## Next time

Recap: typical sets

Formal statement of source coding theorem

Proof of source coding theorem